



Autonomous Career Operations Engine

End-to-End AI-Agentic Orchestration Driven Natively by Antigravity IDE



Generating a Unified Knowledge Graph

Non-Destructive Multimodal Document Consolidation via Native IDE Parsing



Combating Contextual Bloat in Agentic Design

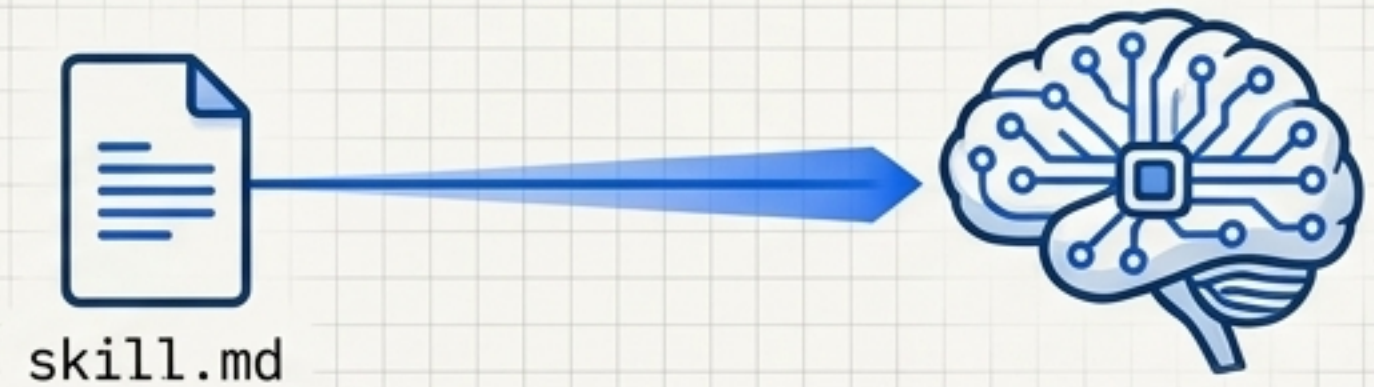
Why Workflows Strictly Utilize .md Skills Over .py Executables

Python Scripts (.py)



Heavy syntax, library imports, and execution logic overhead rapidly consume token limits and severely distract the LLM attention mechanism from the core behavioral task.

Markdown Skills (.md)



Pure Semantic Focus. The agent focuses 100% on behavioral rules and rubrics. Orchestration workflows effortlessly chain multiple markdown skills together without triggering nested, token-heavy script executions.

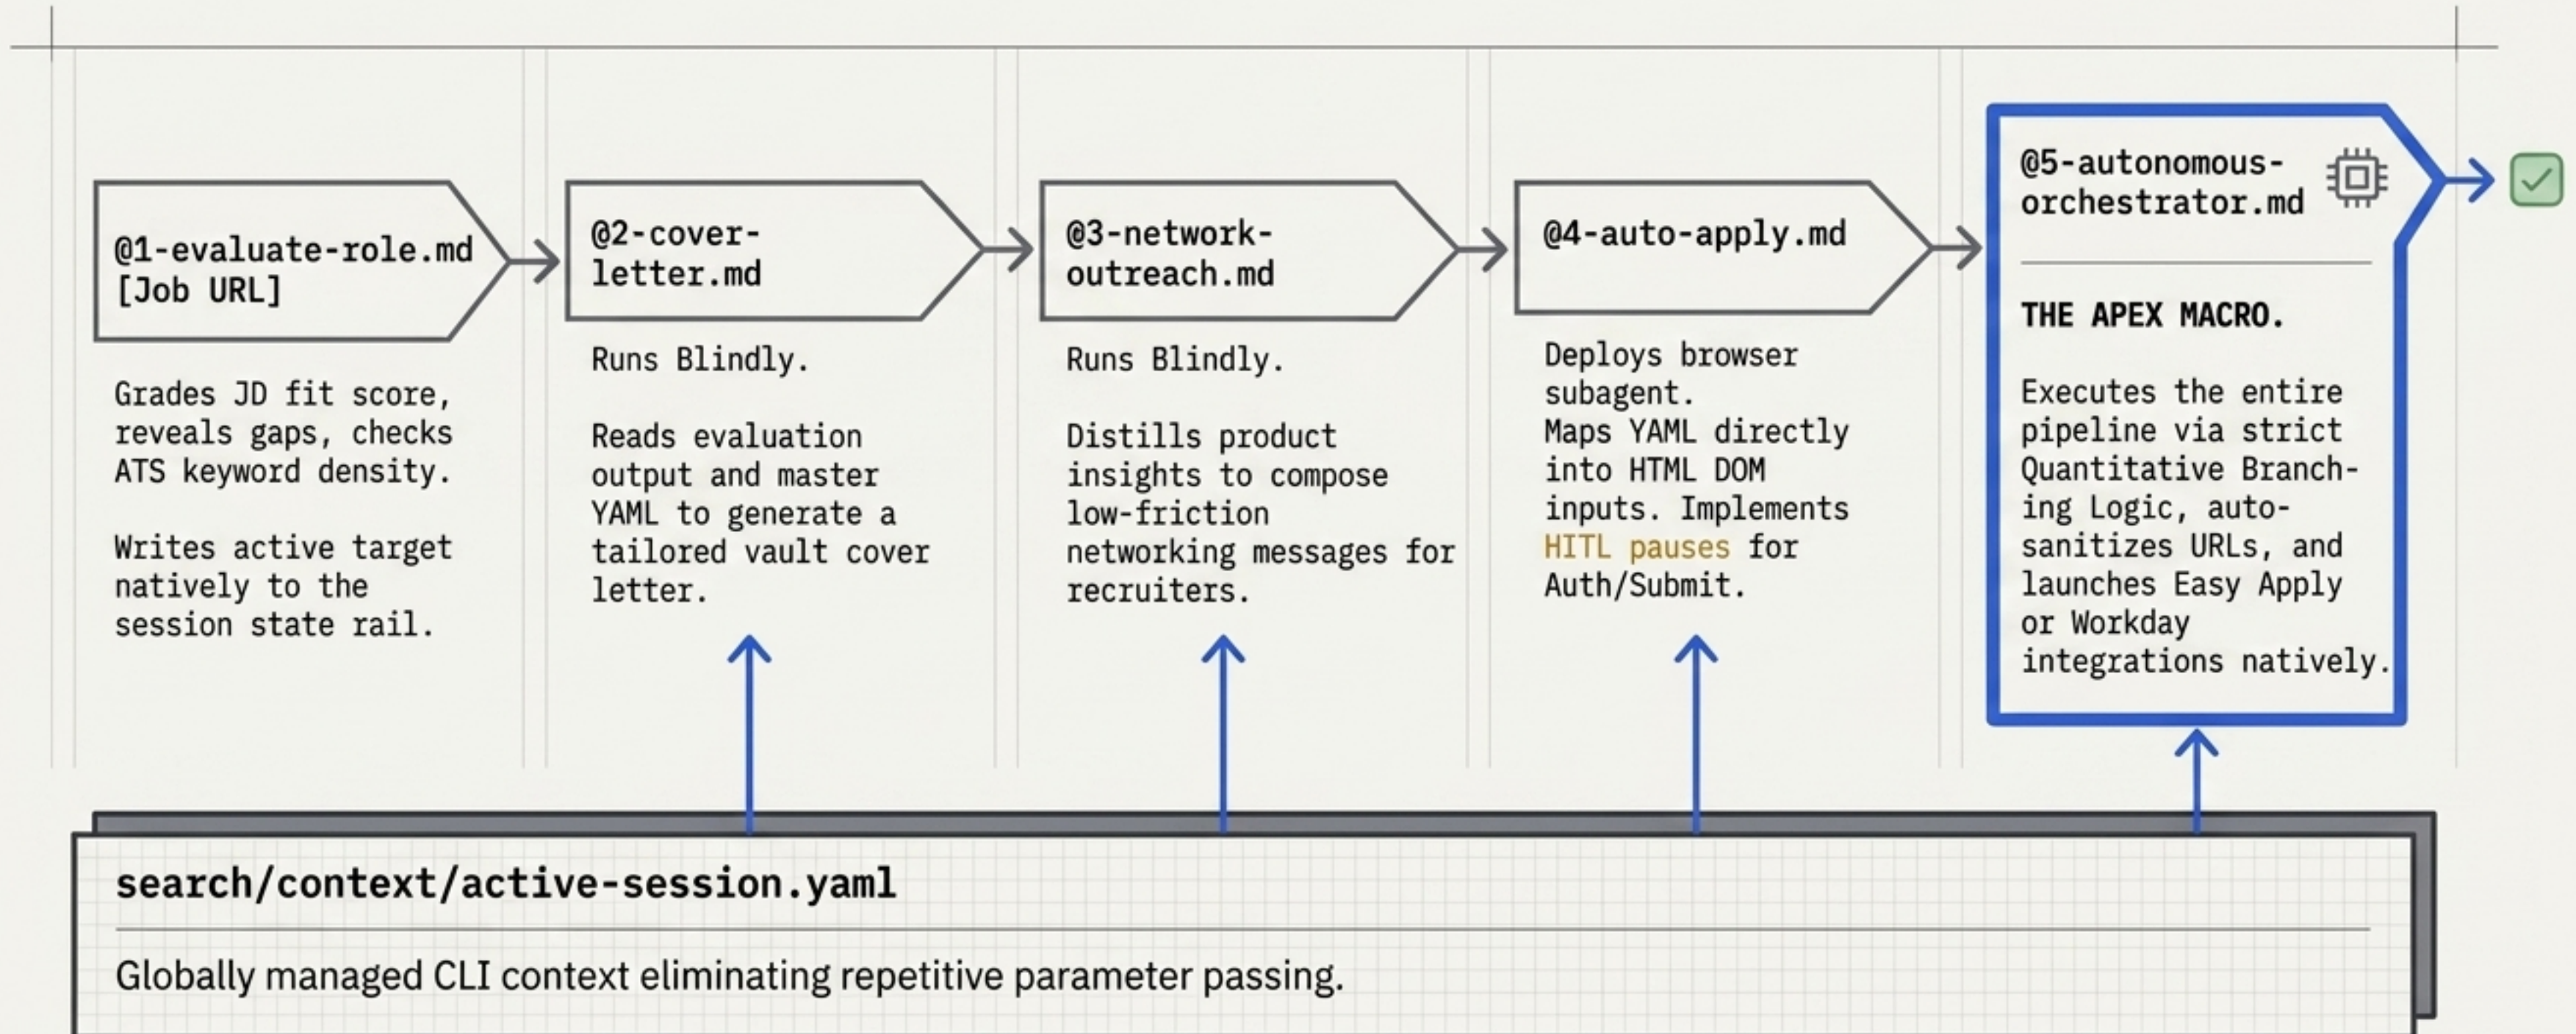


Deterministic Exceptions

Python is isolated strictly for headless absolute sub-shell executions (e.g., URL sanitization via ``urllib.parse`` in ``clean-url/SKILL.md``), invoked via markdown prompts to guarantee zero LLM hallucination.

The Autonomous Macro Pipeline

Deterministic Sequential Execution Driven by Headless State Automation



State-of-the-Art Engineering Breakthroughs

Architectural Milestones in Agentic System Design



The Self-Healing Loop

Continuous Database Evolution. When the browser agent hits unknown proprietary application questions, it pauses via HITL (Human-in-the-Loop, **accented in amber**), asks the user, and natively permanently rewrites the `interview-answers.yaml` database (**in**). It never makes the same mistake twice.



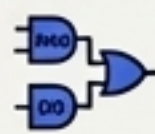
Native Memory Grepping

Zero-Latency Deductions. Bypasses slow AI WebFetch executions. Downstream workflows use explicit `grep_search` **logic** (across the `.md vault`) to instantly map numerical tracking URLs to expanded Company data natively in milliseconds.



Deterministic Anti-Blocking

100% Interception. Proactive sub-shell intercepts strip tracking hooks natively. Mandatory HITL pre-passes (**accented in amber**) pause the subagent to collect manual LinkedIn login states before parsing modal DOMs (**in**).



Multi-Level Logic Gating

Strict Quantitative Barriers. The macro refuses to blindly apply. It halts programmatic browser execution unless three disparate expert skills (`score-jd`, `ats-check`, `recruiter-review`) return a >90% absolute calculation (**in soft forest green**), while **context-switching** between Easy Apply popups and iframes.